

# Linux pipe and subshell

@k1ic

chenkai116@gmail.com



# 管道方式

```
:/home/chenkai3/tmp$ cat get_file_via_pipe.sh  
#!/bin/bash
```

```
function get_content_via_pipe() {  
    res_str='';  
    file_path='./data.txt';  
    cat ${file_path} | while read line  
    do  
        res_str=${res_str}[br]${line};  
        echo ${res_str};  
    done;  
    echo "The Result: "${res_str};  
}
```

```
get_content_via_pipe
```

```
:/home/chenkai3/tmp$ sh get_file_via_pipe.sh
```

```
[br]qqq
```

```
[br]qqq[br]aaa
```

```
[br]qqq[br]aaa[br]zzz
```

```
The Result:
```

# 重定向方式

```
~/home/chenkai3/tmp$ cat get_file_via_redirect.sh
#!/bin/bash

function get_file_via_redirect() {
    res_str='';
    file_path='./data.txt';

    while read line
    do
        res_str=${res_str}[br]${line};
        echo ${res_str};
    done < ${file_path};
    echo "The Result: "${res_str};
}
```

get\_file\_via\_redirect

```
~/home/chenkai3/tmp$ sh get_file_via_redirect.sh
[br]qqq
[br]qqq[br]aaa
[br]qqq[br]aaa[br]zzz
The Result: [br]qqq[br]aaa[br]zzz
```



# 系统调用对比

```
1 Process 2153 attached - interrupt to quit
2 wait4(-1, [{WIFEXITED(s) && WEXITSTATUS(s) == 0}],
  0, NULL) = 2154
3 rt_sigprocmask(SIG_SETMASK, [], NULL, 8) = 0
4 --- SIGCHLD (Child exited) @ 0 (0) ---
5 wait4(-1, 0x7ffffb193fd4, WNOHANG, NULL) = -1 ECHILD
  (No child processes)
6 rt_sigreturn(0xffffffffffffffff) = 0
7 rt_sigaction(SIGINT, {SIG_DFL, [], SA_RESTORER, 0x380fe302d0},
  {0x436c60, [], SA_RESTORER, 0x380fe302d0}, 8) = 0
8 rt_sigprocmask(SIG_BLOCK, NULL, [], 8) = 0
9 read(255, "get_file_via_pipe\n", 273) = 18
10 rt_sigprocmask(SIG_BLOCK, NULL, [], 8) = 0
11 rt_sigprocmask(SIG_BLOCK, [CHLD], [], 8) = 0
12 pipe([3, 4]) = 0
13 rt_sigprocmask(SIG_BLOCK, [INT CHLD], [CHLD], 8) = 0
14 clone(child_stack=0, flags=CLONE_CHILD_CLEARTID|CLONE_CHILD_SETTID|SIGCHLD,
  child_tidptr=0x2b1bcb347e60) = 2162
15 rt_sigprocmask(SIG_SETMASK, [CHLD], NULL, 8) = 0
16 close(4) = 0
17 close(4) = -1 EBADF (Bad file descriptor)
18 rt_sigprocmask(SIG_BLOCK, [INT CHLD], [CHLD], 8) = 0
19 clone(child_stack=0, flags=CLONE_CHILD_CLEARTID|CLONE_CHILD_SETTID|SIGCHLD,
  child_tidptr=0x2b1bcb347e60) = 2163
20 rt_sigprocmask(SIG_SETMASK, [CHLD], NULL, 8) = 0
21 close(3) = 0
22 rt_sigprocmask(SIG_BLOCK, [CHLD], [CHLD], 8) = 0
```

```
1 Process 3694 attached - interrupt to quit
2 wait4(-1, [{WIFEXITED(s) && WEXITSTATUS(s) == 0}],
  0, NULL) = 3695
3 rt_sigprocmask(SIG_SETMASK, [], NULL, 8) = 0
4 --- SIGCHLD (Child exited) @ 0 (0) ---
5 wait4(-1, 0x7ffff8960a224, WNOHANG, NULL) = -1 ECHILD
  (No child processes)
6 rt_sigreturn(0xffffffffffffffff) = 0
7 rt_sigaction(SIGINT, {SIG_DFL, [], SA_RESTORER, 0x380fe302d0},
  {0x436c60, [], SA_RESTORER, 0x380fe302d0}, 8) = 0
8 rt_sigprocmask(SIG_BLOCK, NULL, [], 8) = 0
9 read(255, "get_file_via_redirect\n", 277) = 22
10 rt_sigprocmask(SIG_BLOCK, NULL, [], 8) = 0
11 open("./data.txt", O_RDONLY) = 3
12 fcntl(0, F_GETFD) = 0
13 fcntl(0, F_DUPFD, 10) = 10
14 fcntl(0, F_GETFD) = 0
15 fcntl(10, F_SETFD, FD_CLOEXEC) = 0
16 dup2(3, 0) = 0
17 close(3) = 0
18 ioctl(0, SNDCTL_TMR_TIMEBASE or TCGETS, 0x7ffff89609df0) = -1 ENOTTY
  (Inappropriate ioctl for device)
19 lseek(0, 0, SEEK_CUR) = 0
20 read(0, "qqq\naaa\nzzz\n", 128) = 12
21 lseek(0, -8, SEEK_CUR) = 4
22 rt_sigprocmask(SIG_BLOCK, [CHLD], [], 8) = 0
23 rt_sigprocmask(SIG_SETMASK, [], NULL, 8) = 0
24 open(".", O_RDONLY|O_NONBLOCK|O_DIRECTORY) = 3
25 fcntl(3, F_SETFD, FD_CLOEXEC) = 0
26 getdents(3, /* 9 entries */, 32768) = 296
27 getdents(3, /* 0 entries */, 32768) = 0
28 close(3) = 0
```



VFS inode

# VFS inode

- Sector & Block

$$1\text{Block} = 8\text{Sector} = 8 * 0.5\text{K} = 4\text{K}$$

(Linux OS Default Memory PageSize = 4K)

- File = Blocks + inode

- Inode = file size + uid + gid + rwx permission  
+ time stamp + link total + block Loc



# VFS inode

```
chenkai3@chenkai3:~/tmp$ ll -ai
total 12
1527658 drwxr-xr-x 2 root    root    4096 Mar 25 15:25 .
1527545 drwx----- 9 jianxin jianxin 4096 Mar 25 14:58 ..
 419875 -rw-r--r-- 1 root    root      4 Mar 25 15:25 test.txt
chenkai3@chenkai3:~/tmp$ cat test.txt
123
chenkai3@chenkai3:~/tmp$ stat test.txt
  File: `test.txt'
  Size: 4          Blocks: 8          IO Block: 4096   regular file
Device: 301h/769d Inode: 419875       Links: 1
Access: (0644/-rw-r--r--)  Uid: (  0/   root)   Gid: (  0/   root)
Access: 2015-03-25 15:25:39.000000000 +0800
Modify: 2015-03-25 15:25:19.000000000 +0800
Change: 2015-03-25 15:25:19.000000000 +0800
chenkai3@chenkai3:~/tmp$ ln test.txt test_ln.txt && ln -s test.txt test_s.txt
chenkai3@chenkai3:~/tmp$ stat test.txt
  File: `test.txt'
  Size: 4          Blocks: 8          IO Block: 4096   regular file
Device: 301h/769d Inode: 419875       Links: 2
Access: (0644/-rw-r--r--)  Uid: (  0/   root)   Gid: (  0/   root)
Access: 2015-03-25 15:25:39.000000000 +0800
Modify: 2015-03-25 15:25:19.000000000 +0800
Change: 2015-03-25 15:26:49.000000000 +0800
chenkai3@chenkai3:~/tmp$ ll -ai
total 16
1527658 drwxr-xr-x 2 root    root    4096 Mar 25 15:26 .
1527545 drwx----- 9 jianxin jianxin 4096 Mar 25 14:58 ..
 419875 -rw-r--r-- 2 root    root      4 Mar 25 15:25 test_ln.txt
 419876 lrwxrwxrwx 1 root    root      8 Mar 25 15:26 test_s.txt -> test.txt
 419875 -rw-r--r-- 2 root    root      4 Mar 25 15:25 test.txt
```



# shell子进程

# 子进程

- 概念:

init进程 (`pid=1`) 以后创建的所有进程均称为子进程  
孤儿进程的`ppid=1`

- 创建:

`fork()`创建子进程环境

`exec()`加载、执行进程代码



# 子进程 – 产生场景

- 提交后台作业 &
- 管道 |
- 括号命令列表 ()
- 执行外部脚本

# 进程树

```
:/home/chenkai3/tmp$ pstree -p
init(1)─automount(1695)─{automount}(1696)
                        └─{automount}(1697)
                        └─{automount}(1700)
                        └─{automount}(1703)
      ─cfexecd(18505)
      ─crond(2120)
      ─dnsmasq(21170)
      ─events/0(5)
      ─khelper(6)
      ─krfcommd(10685)
      ─ksoftirqd/0(3)
      ─kthread(15)─aio/0(168)
                  ─ata/0(337)
                  ─ata_aux(338)
                  ─cqueue/0(99)
                  ─kacpid(20)
                  ─kauditd(379)
                  ─kblockd/0(19)
                  ─khubd(102)
                  ─kjournald(354)
                  ─kjournald(1018)
                  ─kmpath_handlerd(994)
                  ─kmpathd/0(993)
                  ─kpsmoused(311)
                  ─kseriod(104)
                  ─kstriped(345)
                  ─kswapd0(167)
                  ─pdflush(9673)
                  ─pdflush(16430)
                  ─rpciod/0(1501)
      ─memcached(1305)─{memcached}(1306)
                      └─{memcached}(1307)
                      └─{memcached}(1308)
```



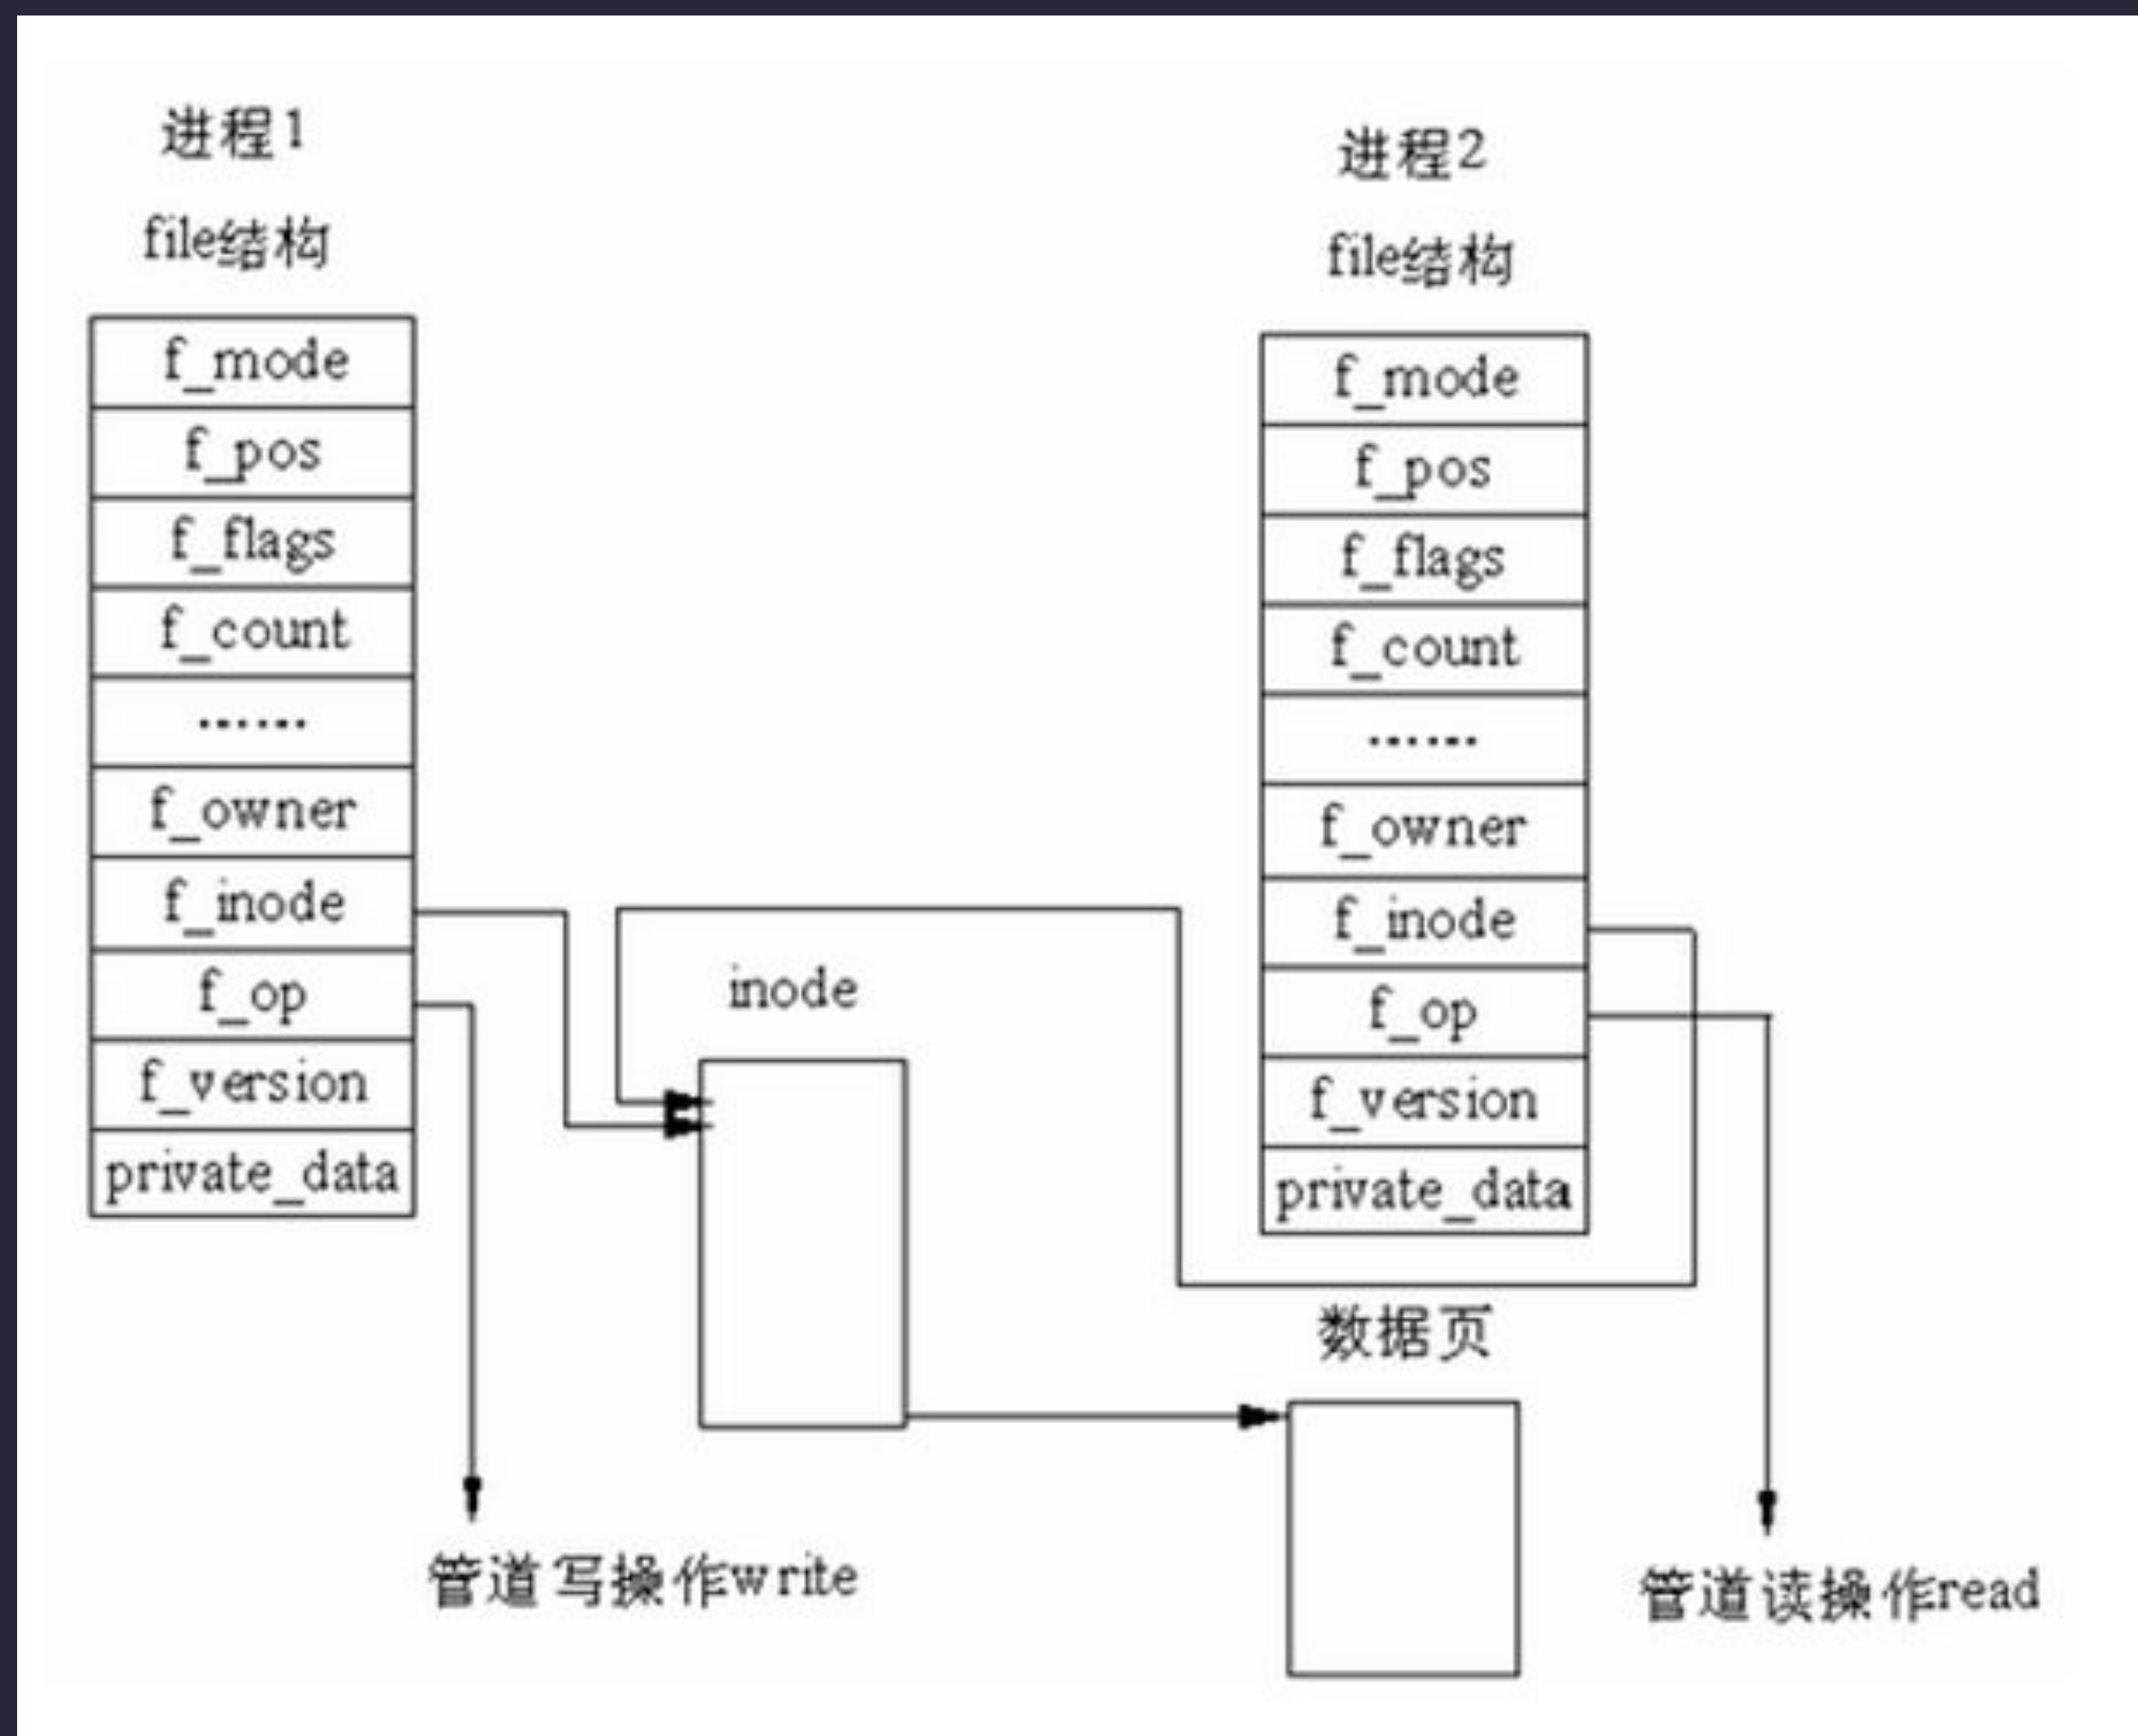


# 管道的特点

- 使用最多的进程间通信方式  
(其他通信机制: 共享内存、消息队列、信号量、套接字等)
- 仅存在于内存, 大小固定  
(一般为4K, 即一个内存页的大小)
- 半双工, 原子读、非原子写  
(管道内数据被读取后管道随即清空, 此时读操作会被阻塞)



# 管道的结构



# 管道的结构

- 两个file结构指向同一个VFS索引节点（inode），索引节点指向一个内存页
- 读入端（fd[0]）、写出端（fd[1]）



# 管道的读写

- 向管道写：将字节流数据复制到inode指向的物理内存  
conditions:

`left memery >= data size && memery unlock`

Yes:

`lock memery --> copy data --> unlock memery`

No:

`write_process sleep in inode's waiting queue until  
all above conditions are met`

# 管道的读写

- 从管道读：一次性复制物理内存中的全部字节流数据
- 读进程可在管道无数据或内存被锁定时立即返回（依赖文件或管道的打开模式），或休眠在inode的等待队列，直到内存被解锁



# 结论

- 管道操作符“|”，将I/O流重定向到一个子进程中，子进程中的变量对于子进程之外的代码块不可见
- 父进程不能访问这些变量，也不能通过`export`方式将子进程变量导出到父进程

Q&A



*THE END*